

01_data_cleaning

May 25, 2025

```
[2]: import pandas as pd
```

```
purchase_df = pd.read_csv('../data/QVI_purchase_behaviour.csv')
transaction_df = pd.read_excel('../data/QVI_transaction_data.xlsx',
                               sheet_name='in')
```

```
[ ]: # Converting 'DATE' column from Excel serial to datetime
transaction_df['DATE'] = pd.to_datetime(transaction_df['DATE'],
                                         origin='1899-12-30', unit='D')

# Basic checks
print("Date Range:", transaction_df['DATE'].min(), "to", transaction_df['DATE'].
      max())
print("Missing Values:\n", transaction_df.isnull().sum())
```

Date Range: 2018-07-01 00:00:00 to 2019-06-30 00:00:00

Missing Values:

DATE	0
STORE_NBR	0
LYLTY_CARD_NBR	0
TXN_ID	0
PROD_NBR	0
PROD_NAME	0
PROD_QTY	0
TOT_SALES	0

dtype: int64

```
[4]: import numpy as np
```

```
# Extracting Pack Size
transaction_df['PACK_SIZE'] = transaction_df['PROD_NAME'].str.
    extract(r'(\d+)g').astype(float)

# Extracting Brand
transaction_df['BRAND'] = transaction_df['PROD_NAME'].str.split().str[0]

# Quick summary
```

```
print(transaction_df[['PROD_NAME', 'PACK_SIZE', 'BRAND']].head())

# Check
print("Missing PACK_SIZE values:", transaction_df['PACK_SIZE'].isna().sum())
```

	PROD_NAME	PACK_SIZE	BRAND
0	Natural Chip Compny SeaSalt	175g	175.0 Natural
1	CCs Nacho Cheese	175g	175.0 CCs
2	Smiths Crinkle Cut Chips Chicken	170g	170.0 Smiths
3	Smiths Chip Thinly S/Cream&Onion	175g	175.0 Smiths
4	Kettle Tortilla ChpsHny&Jlpno Chili	150g	150.0 Kettle

Missing PACK_SIZE values: 6064

```
[5]: # Show rows with missing PACK_SIZE
missing_pack = transaction_df[transaction_df['PACK_SIZE'].isna()]
print(missing_pack[['PROD_NAME']].head(10))
```

	PROD_NAME
9	Grain Waves Sour Cream&Chives 210G
34	Red Rock Deli Sp Salt & Truffle 150G
35	Smiths Thinly Swt Chli&S/Cream 175G
212	Red Rock Deli Sp Salt & Truffle 150G
292	Grain Waves Sour Cream&Chives 210G
331	Red Rock Deli Sp Salt & Truffle 150G
421	Smiths Thinly Swt Chli&S/Cream 175G
423	Grain Waves Sour Cream&Chives 210G
428	Smiths Thinly Swt Chli&S/Cream 175G
465	Red Rock Deli Sp Salt & Truffle 150G

```
[6]: # Update PACK_SIZE extraction to include both 'g' and 'G'
transaction_df['PACK_SIZE'] = transaction_df['PROD_NAME'].str.
    ↪extract(r'(\d+)[gG]').astype(float)

# Re-check missing values
print("Missing PACK_SIZE values after fix:", transaction_df['PACK_SIZE'].isna().
    ↪sum())
```

Missing PACK_SIZE values after fix: 0

```
[7]: duplicates = transaction_df.duplicated().sum()
print(f"Duplicate rows: {duplicates}")

print(transaction_df['PROD_QTY'].describe())

# View extreme values
print(transaction_df[transaction_df['PROD_QTY'] > 5].sort_values(by='PROD_QTY',
    ↪ascending=False).head(10))
```

Duplicate rows: 1

```
count    264836.000000
mean      1.907309
std       0.643654
min       1.000000
25%       2.000000
50%       2.000000
75%       2.000000
max       200.000000
```

Name: PROD_QTY, dtype: float64

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
69762	2018-08-19	226	226000	226201	4	
69763	2019-05-20	226	226000	226210	4	

	PROD_NAME	PROD_QTY	TOT_SALES	PACK_SIZE	\
69762	Dorito Corn Chp	Supreme 380g	200	650.0	380.0
69763	Dorito Corn Chp	Supreme 380g	200	650.0	380.0

	BRAND
69762	Dorito
69763	Dorito

```
[8]: # Drop duplicate row
transaction_df.drop_duplicates(inplace=True)

# Filter out unrealistic bulk buys (e.g., anything over 5 units of chips)
transaction_df = transaction_df[transaction_df['PROD_QTY'] <= 5]

# Confirm cleanup
print("Remaining rows after cleaning:", len(transaction_df))
```

Remaining rows after cleaning: 264833

```
[10]: customer_df = pd.read_csv('../data/QVI_purchase_behaviour.csv')
```

```
[12]: import pandas as pd

customer_df = pd.read_csv('../data/QVI_purchase_behaviour.csv')

# Preview the data
print("Customer Data Preview:")
print(customer_df.head())

# Check for missing values
print("\nMissing values in customer data:")
print(customer_df.isnull().sum())
```

Customer Data Preview:

	LYLTY_CARD_NBR	LIFESTAGE	PREMIUM_CUSTOMER
0	1000	YOUNG SINGLES/COUPLES	Premium
1	1002	YOUNG SINGLES/COUPLES	Mainstream
2	1003	YOUNG FAMILIES	Budget
3	1004	OLDER SINGLES/COUPLES	Mainstream
4	1005	MIDAGE SINGLES/COUPLES	Mainstream

Missing values in customer data:

```
LYLTY_CARD_NBR    0
LIFESTAGE         0
PREMIUM_CUSTOMER  0
dtype: int64
```

```
[13]: merged_df = pd.merge(transaction_df, customer_df, on='LYLTY_CARD_NBR',
    ↪how='inner')

# Quick check
print("Merged Data Preview:")
print(merged_df.head())

# Shape and null check
print("\nShape of merged data:", merged_df.shape)
print("Missing values after merge:")
print(merged_df.isnull().sum())
```

Merged Data Preview:

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	2018-10-17	1	1000	1	5	
1	2019-05-14	1	1307	348	66	
2	2019-05-20	1	1343	383	61	
3	2018-08-17	2	2373	974	69	
4	2018-08-18	2	2426	1038	108	

	PROD_NAME	PROD_QTY	TOT_SALES	PACK_SIZE	\
0	Natural Chip Compny SeaSalt175g	2	6.0	175.0	
1	CCs Nacho Cheese 175g	3	6.3	175.0	
2	Smiths Crinkle Cut Chips Chicken 170g	2	2.9	170.0	
3	Smiths Chip Thinly S/Cream&Union 175g	5	15.0	175.0	
4	Kettle Tortilla ChpsHny&Jlpno Chili 150g	3	13.8	150.0	

	BRAND	LIFESTAGE	PREMIUM_CUSTOMER
0	Natural	YOUNG SINGLES/COUPLES	Premium
1	CCs	MIDAGE SINGLES/COUPLES	Budget
2	Smiths	MIDAGE SINGLES/COUPLES	Budget
3	Smiths	MIDAGE SINGLES/COUPLES	Budget
4	Kettle	MIDAGE SINGLES/COUPLES	Budget

Shape of merged data: (264833, 12)

Missing values after merge:

DATE	0
STORE_NBR	0
LYLTY_CARD_NBR	0
TXN_ID	0
PROD_NBR	0
PROD_NAME	0
PROD_QTY	0
TOT_SALES	0
PACK_SIZE	0
BRAND	0
LIFESTAGE	0
PREMIUM_CUSTOMER	0

dtype: int64

```
[14]: import matplotlib.pyplot as plt
import seaborn as sns

# Total sales by LIFESTAGE and PREMIUM_CUSTOMER
sales_by_segment = merged_df.groupby(['LIFESTAGE',
    ↪ 'PREMIUM_CUSTOMER'])['TOT_SALES'].sum().reset_index()



plt.figure(figsize=(12,6))
sns.barplot(data=sales_by_segment, x='LIFESTAGE', y='TOT_SALES',
    ↪ hue='PREMIUM_CUSTOMER')
plt.title('Total Sales by Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

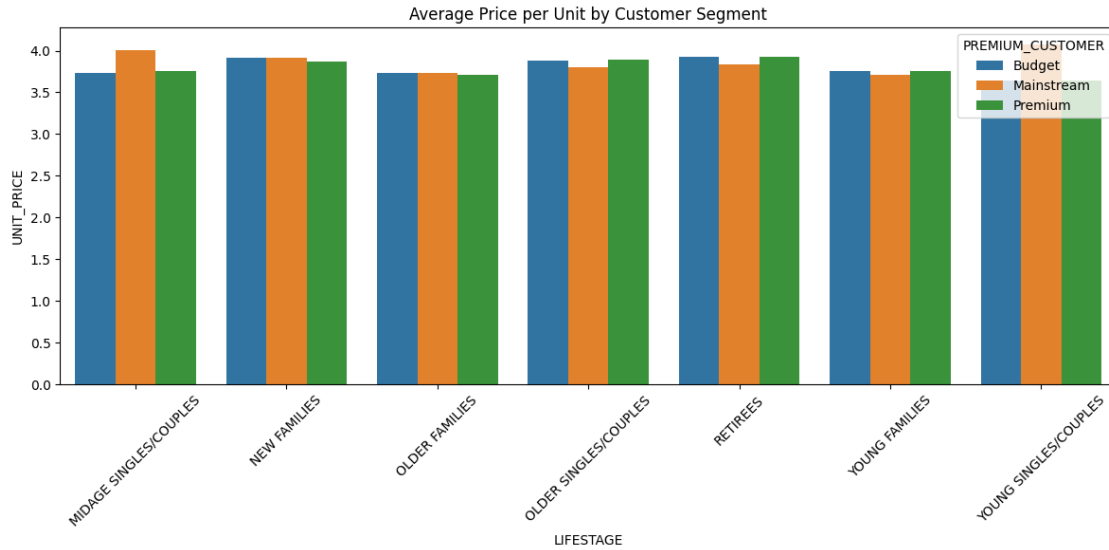
# Units sold by segment
units_by_segment = merged_df.groupby(['LIFESTAGE',
    ↪ 'PREMIUM_CUSTOMER'])['PROD_QTY'].sum().reset_index()

plt.figure(figsize=(12,6))
sns.barplot(data=units_by_segment, x='LIFESTAGE', y='PROD_QTY',
    ↪ hue='PREMIUM_CUSTOMER')
plt.title('Total Units Sold by Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Average price per unit by segment
merged_df['UNIT_PRICE'] = merged_df['TOT_SALES'] / merged_df['PROD_QTY']
avg_price_by_segment = merged_df.groupby(['LIFESTAGE',
    ↪ 'PREMIUM_CUSTOMER'])['UNIT_PRICE'].mean().reset_index()
```

```
plt.figure(figsize=(12,6))
sns.barplot(data=avg_price_by_segment, x='LIFESTAGE', y='UNIT_PRICE',
            hue='PREMIUM_CUSTOMER')
plt.title('Average Price per Unit by Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```





```
[16]: merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 264833 entries, 0 to 264832
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DATE                   264833 non-null  datetime64[ns]
1   STORE_NBR              264833 non-null  int64
2   LYLTY_CARD_NBR         264833 non-null  int64
3   TXN_ID                 264833 non-null  int64
4   PROD_NBR               264833 non-null  int64
5   PROD_NAME              264833 non-null  object
6   PROD_QTY               264833 non-null  int64
7   TOT_SALES              264833 non-null  float64
8   PACK_SIZE              264833 non-null  float64
9   BRAND                  264833 non-null  object
10  LIFESTAGE               264833 non-null  object
11  PREMIUM_CUSTOMER        264833 non-null  object
12  UNIT_PRICE              264833 non-null  float64
dtypes: datetime64[ns](1), float64(3), int64(5), object(4)
memory usage: 26.3+ MB
```

```
[17]: merged_df.describe()
```

```
[17]:
```

	DATE	STORE_NBR	LYLTY_CARD_NBR
count	264833	264833.000000	2.648330e+05
mean	2018-12-30 00:52:39.666657792	135.079529	1.355489e+05
min	2018-07-01 00:00:00	1.000000	1.000000e+03

25%	2018-09-30 00:00:00	70.000000	7.002100e+04
50%	2018-12-30 00:00:00	130.000000	1.303570e+05
75%	2019-03-31 00:00:00	203.000000	2.030940e+05
max	2019-06-30 00:00:00	272.000000	2.373711e+06
std	NaN	76.784189	8.058003e+04

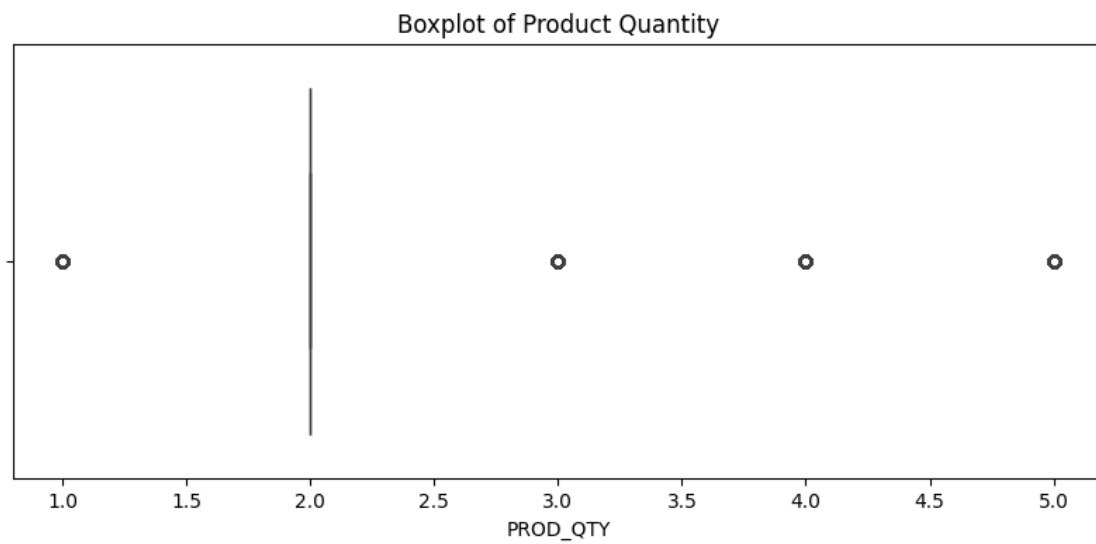

	TXN_ID	PROD_NBR	PROD_QTY	TOT_SALES \
count	2.648330e+05	264833.000000	264833.000000	264833.000000
mean	1.351577e+05	56.583598	1.905812	7.299351
min	1.000000e+00	1.000000	1.000000	1.500000
25%	6.760000e+04	28.000000	2.000000	5.400000
50%	1.351370e+05	56.000000	2.000000	7.400000
75%	2.027000e+05	85.000000	2.000000	9.200000
max	2.415841e+06	114.000000	5.000000	29.500000
std	7.813305e+04	32.826498	0.343437	2.527244

	PACK_SIZE	UNIT_PRICE
count	264833.000000	264833.000000
mean	182.425540	3.824632
min	70.000000	1.320000
25%	150.000000	3.000000
50%	170.000000	3.800000
75%	175.000000	4.600000
max	380.000000	6.500000
std	64.325268	1.109527


```
[19]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 4))
sns.boxplot(x=merged_df['TOT_SALES'])
plt.title('Boxplot of Total Sales')
plt.show()

plt.figure(figsize=(10, 4))
sns.boxplot(x=merged_df['PROD_QTY'])
plt.title('Boxplot of Product Quantity')
plt.show()
```

```
[20]: plt.figure(figsize=(10, 4))
sns.boxplot(x=merged_df['UNIT_PRICE'])
plt.title('Boxplot of Unit Price')
plt.show()
```



```
[21]: # Check high quantity purchases
merged_df['PROD_QTY'].value_counts().sort_index(ascending=False).head(10)
```

```
[21]: PROD_QTY
5      450
4      397
3      430
2    236038
1    27518
Name: count, dtype: int64
```

```
[22]: import matplotlib.pyplot as plt
import seaborn as sns

# Set plot style
sns.set(style="whitegrid")

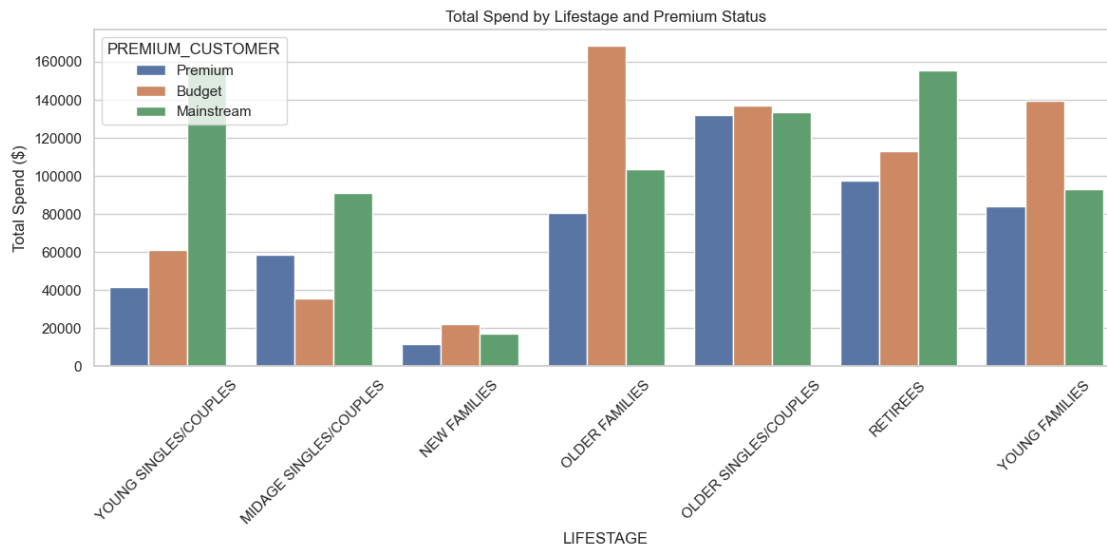
# Grouped spend by LIFESTAGE and PREMIUM_CUSTOMER
plt.figure(figsize=(12, 6))
sns.barplot(data=merged_df,
            x='LIFESTAGE',
            y='TOT_SALES',
            hue='PREMIUM_CUSTOMER',
            estimator=sum,
            ci=None)
plt.title('Total Spend by Lifestage and Premium Status')
plt.xticks(rotation=45)
plt.ylabel('Total Spend ($)')
plt.tight_layout()
```

```
plt.show()
```


C:\Users\Arnav\AppData\Local\Temp\ipykernel_5852\3628436135.py:9: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

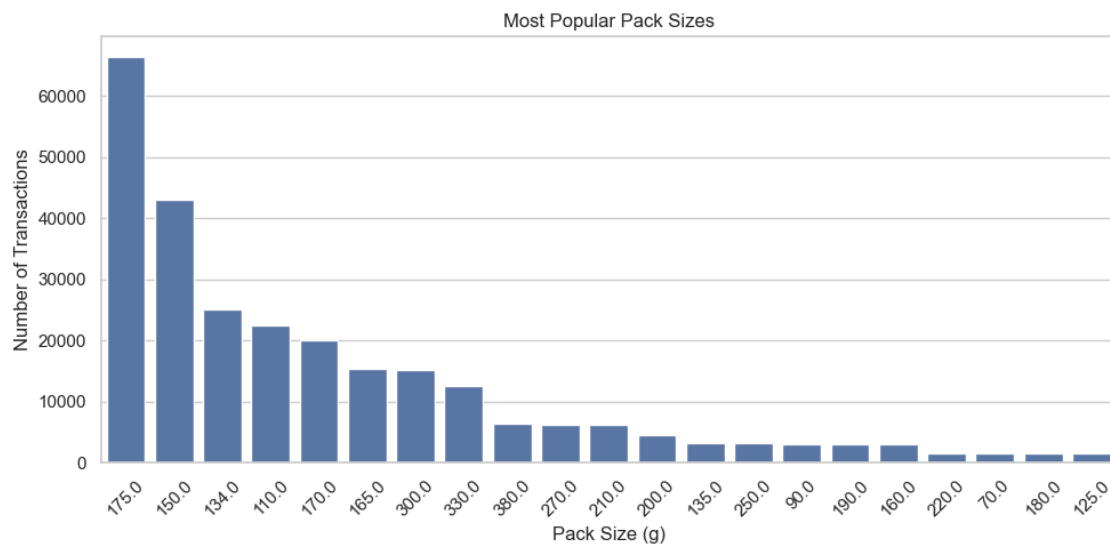
```
sns.barplot(data=merged_df,
```



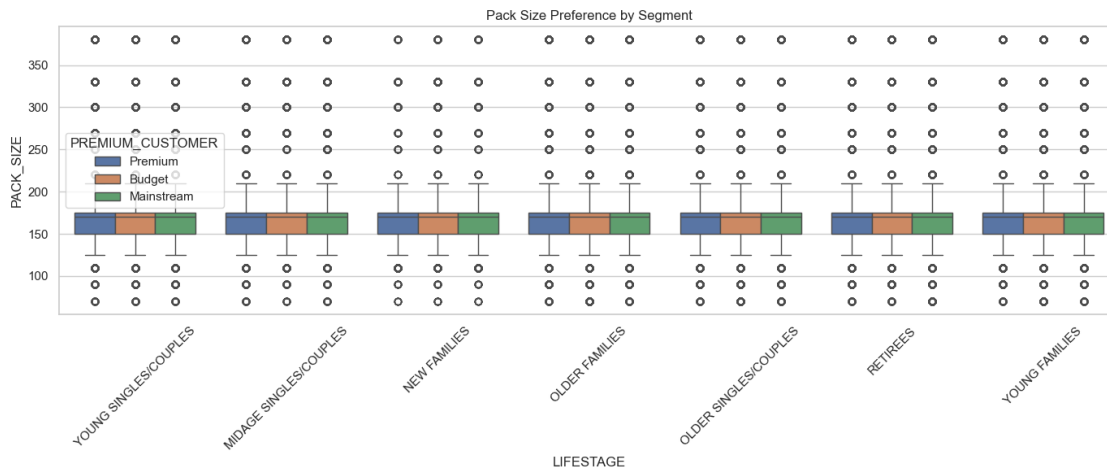
```
[23]: # Average spend by segment
plt.figure(figsize=(12, 6))
sns.barplot(data=merged_df,
            x='LIFESTAGE',
            y='TOT_SALES',
            hue='PREMIUM_CUSTOMER',
            estimator='mean',
            errorbar=None)
plt.title('Average Spend per Transaction by Segment')
plt.xticks(rotation=45)
plt.ylabel('Average Spend ($)')
plt.tight_layout()
plt.show()
```



```
[24]: plt.figure(figsize=(10, 5))
sns.countplot(data=merged_df, x='PACK_SIZE', order=merged_df['PACK_SIZE'].
↳value_counts().index)
plt.title('Most Popular Pack Sizes')
plt.xticks(rotation=45)
plt.xlabel('Pack Size (g)')
plt.ylabel('Number of Transactions')
plt.tight_layout()
plt.show()
```

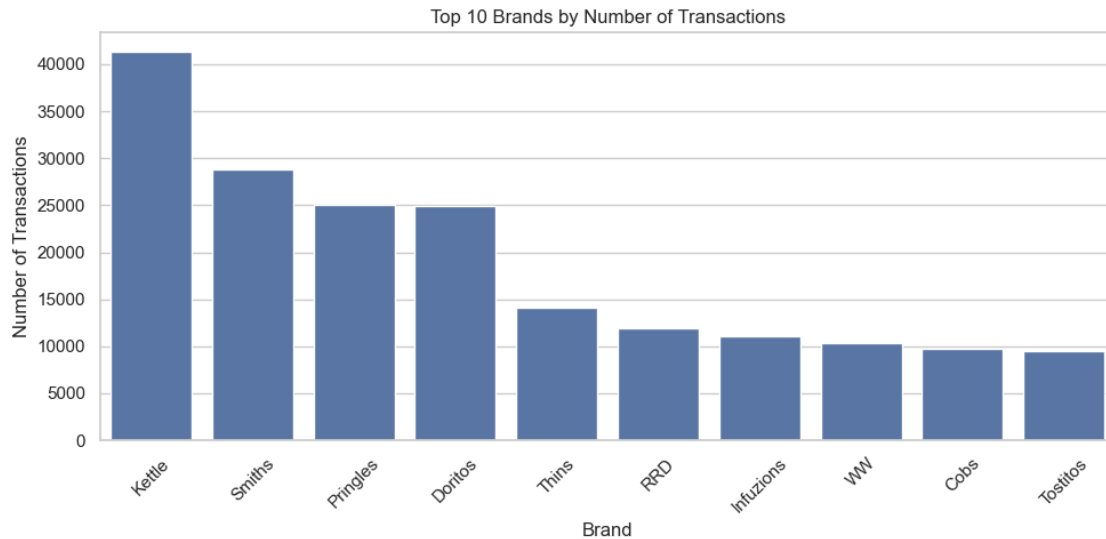


```
[25]: plt.figure(figsize=(14, 6))
sns.boxplot(data=merged_df, x='LIFESTAGE', y='PACK_SIZE',
            hue='PREMIUM_CUSTOMER')
plt.title('Pack Size Preference by Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

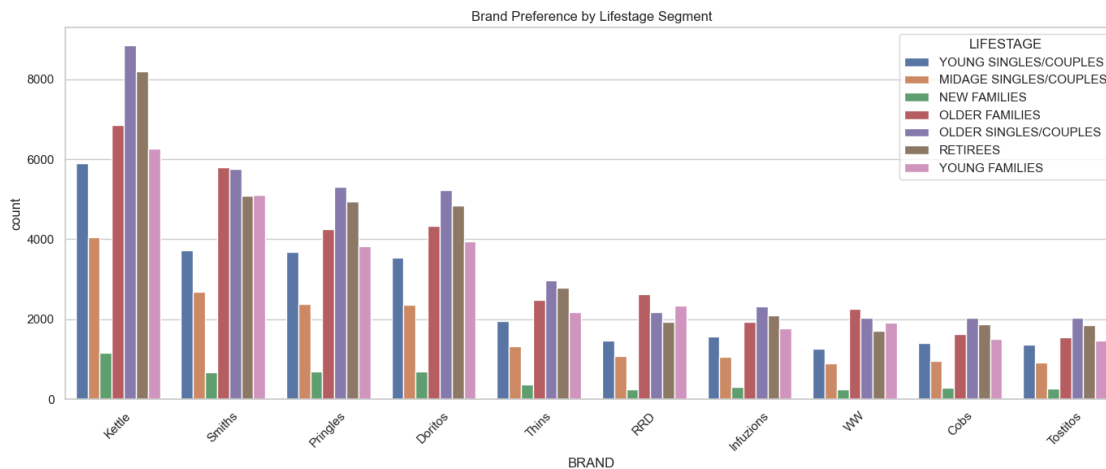


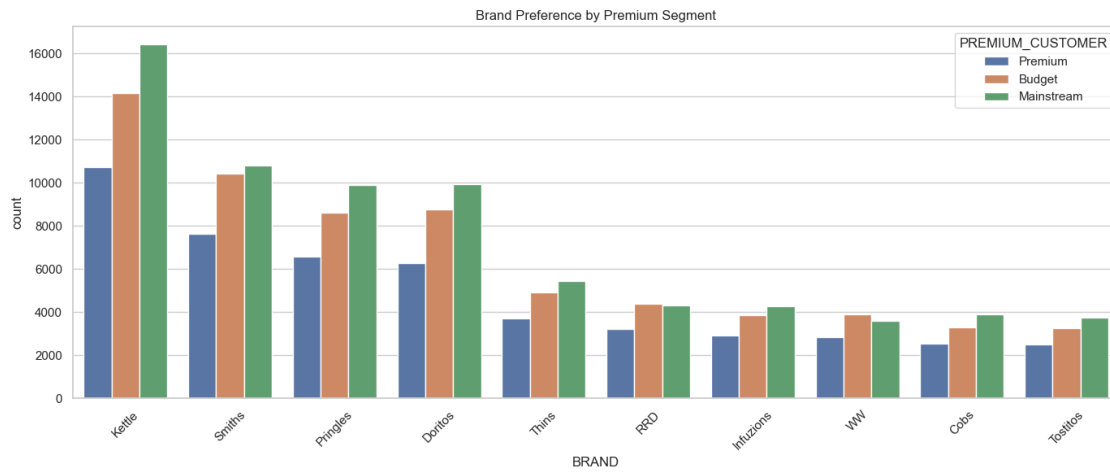
```
[26]: top_brands = merged_df['BRAND'].value_counts().head(10)

plt.figure(figsize=(10, 5))
sns.barplot(x=top_brands.index, y=top_brands.values)
plt.title('Top 10 Brands by Number of Transactions')
plt.ylabel('Number of Transactions')
plt.xlabel('Brand')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



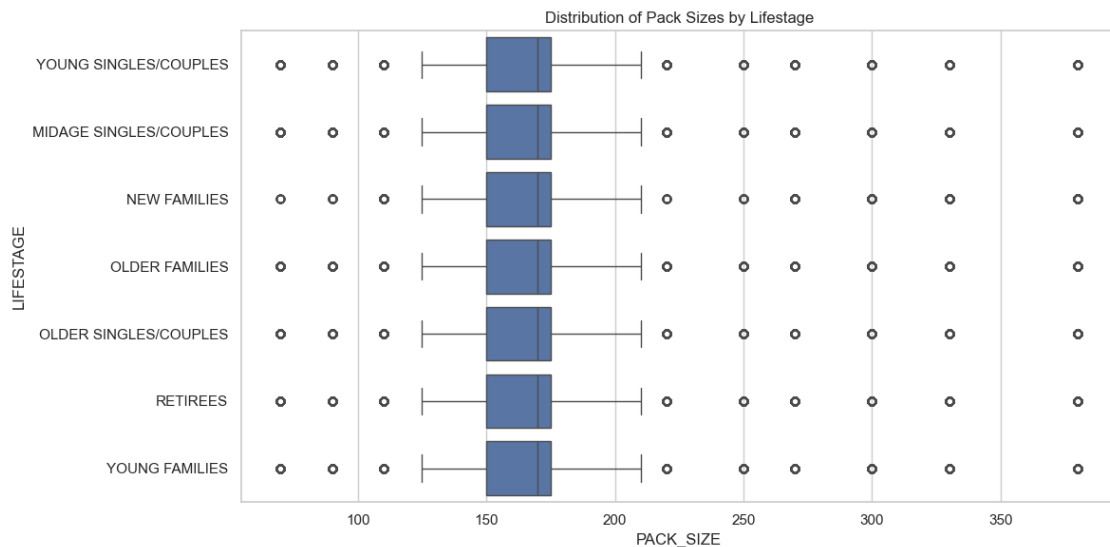
```
[27]: plt.figure(figsize=(14, 6))
sns.countplot(data=merged_df, x='BRAND', hue='LIFESTAGE', order=top_brands.
↪index)
plt.title('Brand Preference by Lifestage Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
plt.figure(figsize=(14, 6))
sns.countplot(data=merged_df, x='BRAND', hue='PREMIUM_CUSTOMER',
↪order=top_brands.index)
plt.title('Brand Preference by Premium Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

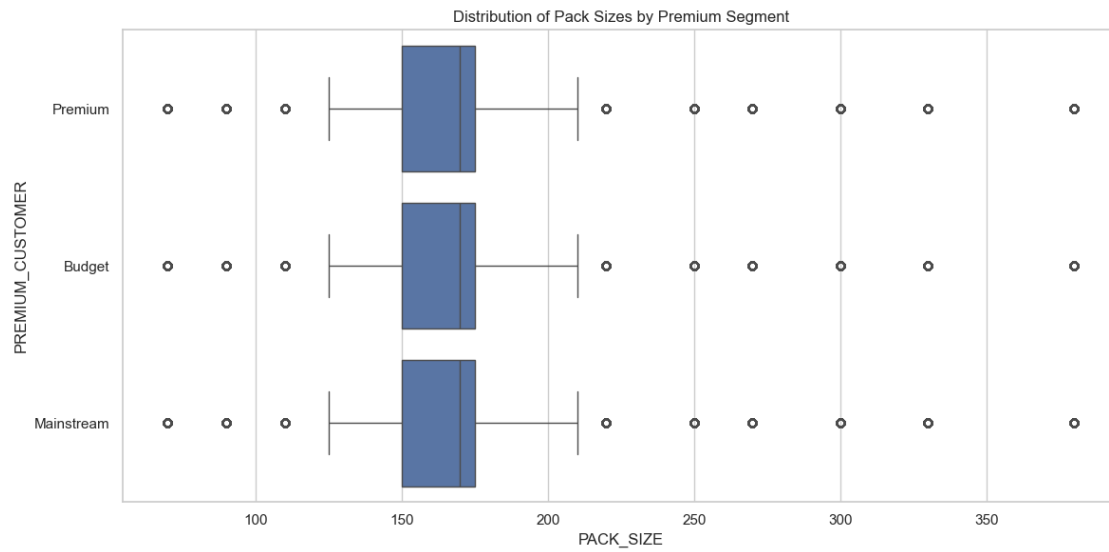




```
[28]: plt.figure(figsize=(12, 6))
sns.boxplot(data=merged_df, x='PACK_SIZE', y='LIFESTAGE')
plt.title('Distribution of Pack Sizes by Lifestage')
plt.tight_layout()
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(data=merged_df, x='PACK_SIZE', y='PREMIUM_CUSTOMER')
plt.title('Distribution of Pack Sizes by Premium Segment')
plt.tight_layout()
plt.show()
```





[]: